

Data scientist focused on the operational use of analysis: forecasting, experiment interpretation, dashboards, and reviewable workflows. Brings an ML engineering perspective to decision support without removing human approval.

Professional Experience

Northstar Transit Labs (fictional) <i>ML Engineer</i>	2022–present
○ Built a monitored document-classification service with human review and reproducible evaluation datasets. ○ Reduced manual triage time by 35% in a fictional internal pilot while preserving reviewer approval. ○ Built an internal retrieval assistant with source citations and an evaluation set. ○ 35% less manual triage (pilot) Reviewer approval retained	
Meridian Analytics (fictional) <i>Data Scientist</i>	2018–2022
○ Used Python and SQL for forecasting, experiment analysis, and dashboards. ○ Coordinated a small cross-functional working group and mentored junior analysts.	
Example University (fictional) <i>Independent research project</i>	2017–2018
Evaluated uncertainty in image classification on a synthetic dataset and shared reproducible notebooks internally. This was not a peer-reviewed publication.	

Core Expertise

Analytics: Python, SQL, experiment analysis, forecasting, and dashboards.
Decision Support: Cross-functional coordination, reproducible evidence, and human approval.

Relevant Project Evidence

Analysis & Teams: Combined forecasting, experiment analysis, and dashboards with cross-functional coordination at Meridian.
Operational Outcome: Northstar's fictional internal pilot reduced manual triage time by 35%; this was a workflow result, not a model-accuracy claim or company-wide rollout.

Education & Profile

Example University (fictional) <i>BSc in Computer Science</i> Fictional degree for the public demo.	2017
Profile: Fictional portfolio example	