

ML engineer focused on retrieval-based assistants with source citations and evaluation sets. Combines this work with experience in monitored classification and human-review workflows.

Professional Experience

Northstar Transit Labs (fictional)

ML Engineer

2022–present

- Built a monitored document-classification service with human review and reproducible evaluation datasets.
- Reduced manual triage time by 35% in a fictional internal pilot while preserving reviewer approval.
- Built an internal retrieval assistant with source citations and an evaluation set.
- `Source-cited answers` `Dedicated evaluation set`

Meridian Analytics (fictional)

Data Scientist

2018–2022

- Used Python and SQL for forecasting, experiment analysis, and dashboards.
- Coordinated a small cross-functional working group and mentored junior analysts.

Example University (fictional)

Independent research project

2017–2018

Evaluated uncertainty in image classification on a synthetic dataset and shared reproducible notebooks internally. This was not a peer-reviewed publication.

Core Expertise

Applied AI: Retrieval, source citations, evaluation datasets, and model monitoring.

Implementation: Python, FastAPI, React, SQL, and human-review workflows.

Relevant Project Evidence

Grounded Retrieval: Built Northstar’s internal knowledge-base assistant with source citations so answers could be checked against their sources.

Evaluation & Review: Created an evaluation set for the assistant; in the separate classification service, retained human approval of results.

Education & Profile

Example University (fictional)

BSc in Computer Science

2017

Fictional degree for the public demo.

Profile: Fictional portfolio example