

ML engineer with a data-science background, focused on turning models into monitored services. Builds evaluation datasets for classification and retrieval, retaining reviewer approval in the classification service.

Professional Experience

Northstar Transit Labs (fictional) <i>ML Engineer</i>	2022–present
○ Built a monitored document-classification service with human review and reproducible evaluation datasets. ○ Reduced manual triage time by 35% in a fictional internal pilot while preserving reviewer approval. ○ Built an internal retrieval assistant with source citations and an evaluation set. ○ Monitored classification Human approval retained	
Meridian Analytics (fictional) <i>Data Scientist</i>	2018–2022
○ Used Python and SQL for forecasting, experiment analysis, and dashboards. ○ Coordinated a small cross-functional working group and mentored junior analysts.	
Example University (fictional) <i>Independent research project</i>	2017–2018
Evaluated uncertainty in image classification on a synthetic dataset and shared reproducible notebooks internally. This was not a peer-reviewed publication.	

Core Expertise

ML Engineering: Python, PyTorch, FastAPI, model evaluation, and production observability.
Data & Interfaces: SQL, reproducible evaluation datasets, React, and human-review workflows.

Relevant Project Evidence

Service Delivery: At Northstar, combined document classification, monitoring, and reviewer approval in one service.
Evaluation: Built reproducible datasets for the classification service and an evaluation set for the retrieval assistant.

Education & Profile

Example University (fictional) <i>BSc in Computer Science</i> Fictional degree for the public demo.	2017
Profile: Fictional portfolio example	